
Introduction to High-Performance Computing

Exercise/5

Giorgio Amati

Corso di dottorato in Ingegneria Aeronautica e Spaziale 2024

g.amati@cineca.it / g.amaticode@gmail.com

Agenda

✓ Exercise

- OpenMP
- OpenACC
- OpenMP-Offload
- Cuda-Fortran
- Do concurrent
- OpenCL
- Matmul
- MPI
- MPI+OpenMP

Code Structure/1

- ✓ Each directory contains source code and a script for compiling and one script for submitting, plus a README file
 - Source code in Fortran & in C (when possible)

```
|— LESSON_1  (openMP)
|— LESSON_2  (openACC)
|— LESSON_3  (openMP offload)
|— LESSON_4  (CUDA)
|— LESSON_5  (do concurrent)
|— LESSON_6  (OpenCL)
|— LESSON_7  (matmul)
|— LESSON_8  (MPI)
|— LESSON_9  (MPI+OpenMP)
```

Code Structure/2

✓ Example: LESSON_3

```
clean.sh
compile.C.sh
compile.fortran.sh
inc_precision.h
mm.c
mm.F90
mod_tools.F90
README
submit.c.slurm
submit.fortran.slurm
```

Code Structure/3

- ✓ The original code is the same so you can see the differences between different parallel paradigm

```
$ diff LESSON_3/mm.c LESSON_2/mm.c
53c53
<  #pragma omp target teams distribute parallel for
collapse(2) map(to:a,b) map(from:c)
---
>  #pragma acc kernels
55,56c55,56
<      for (j = 0; j < nn; j++) {
<          for (k = 0; k < nn; k++) {
---
>      for (k = 0; k < nn; k++) {
>          for (j = 0; j < nn; j++) {
```

Step to do

- ✓ Clean all (if necessary)
- ✓ Compile
- ✓ launch Job

```
[gamati01@login14 LESSON_3]$ ./clean.sh
cleaning directory....
.... all done
[gamati01@login14 LESSON_3]$ ./compile.C.sh
Currently Loaded Modulefiles:
  1) profile/base    2) nvhpc/21.9
```

Key:

default-version

compiling with nvc -Minfo=acc -O2 -mp=gpu

main:

54, Generating map(to:a[:] [:])

Generating map(from:c[:] [:])

Generating map(to:b[:] [:])

That's all folks!!!

Check results

- ✓ Check reports a single value of the result.
 - Because it is a product on n random-number between 0 and 1 it should be, from a statistical point of view around $\langle n/4 \rangle$

```
=====
double precision
size 2048
=====
Initialization
Elapsed time for initialization
Total time -----> 0.066232
=====
Time -----> 3.200976
Mflops -----> 5367.072163
Check -----> 518.636322
```

Why Matrix-Matrix Multiplication?

- ✓ It is simple and, for size big enough, able to exploit CPU & GPU performance
- ✓ You can easily apply many HPC tricks & features
 - Three nested loops
 - 1 instruction
 - $2 \cdot n^3$ Flops
 - n^2 Memory
 - 3 load + 1 store per iteration

```
do j = 1, n
  do k = 1, n
    do i = 1, n
      c(i,j) = c(i,j) + a(i,k)*b(k,j)
    enddo
  enddo
enddo
```


LESSON_1 (cpu)

- ✓ `nvc -Minfo=mp -O2 -mp`
- ✓ Try to measure performance (Mflops) increasing the number of threads
- ✓ Does the Speed-up depends from the size of the problem?
- ✓ Number of thread controlled via environmental variable
 - `export OMP_NUM_THREADS=2`

```
#pragma omp parallel shared(a,b,c),private(i,j,k)
#pragma omp for
for (i = 0; i < nn; i++) {
    for (k = 0; k < nn; k++) {
        for (j = 0; j < nn; j++) {
            c[i][j] = c[i][j] + a[i][k]*b[k][j];
        }
    }
}
```

...

LESSON_2 (GPU)

- ✓ `nvfortran/nvc -Minfo=acc -O2 -acc`
- ✓ Verify performance for version `mm.0.F90` and version `mm.1.F90`
- ✓ Does performance changes with different size?
- ✓ Does C code present the same behaviour?

```
!$acc data copyin(a,b) copy(c)
!$acc parallel
  do j = 1, n
    do k = 1, n
      do i = 1, n
        c(i,j) = c(i,j) + a(i,k)*b(k,j)
      enddo
    enddo
  enddo
!$acc end parallel
!$acc end data
```

```
!$acc data copyin(a,b) copy(c)
!$acc parallel collapse(2)
  do j = 1, n
    do i = 1, n
!$acc loop seq
      do k = 1, n
        c(i,j) = c(i,j) + a(i,k)*b(k,j)
      enddo
    enddo
  enddo
!$acc end parallel
!$acc end data
```

LESSON_3 (GPU)

- ✓ Does the OMP Offload version behave as OpenACC (in performance)?
 - `nvfortran -Minfo=acc -O2 -mp=gpu`

```
!$OMP target teams distribute parallel do map(to:a,b) map(from:c)
  do j = 1, n
    do i = 1, n
      do k = 1, n
        c(i,j) = c(i,j) + a(i,k)*b(k,j)
      enddo
    enddo
  enddo
```

LESSON_4: CUDA Fortran (GPU)

✓ Again: what about performances?

○ `nvfortran -Minfo=acc -O2 -Mcuda`

```
real(my_kind), dimension(:,:), allocatable:: a ! matrix (origin)
real(my_kind), dimension(:,:), allocatable:: b ! matrix (origin)
real(my_kind), dimension(:,:), allocatable:: c ! matrix
real(my_kind), dimension(:,:), device, allocatable:: a_gpu ! matrix (origin)
real(my_kind), dimension(:,:), device, allocatable:: b_gpu ! matrix (origin)
real(my_kind), dimension(:,:), device, allocatable:: c_gpu ! matrix

a_gpu = a
b_gpu = b
c_gpu = c
!$cuf kernel do(2) <<<*,*>>>
do j = 1, n
  do j = 1, n
    do k = 1, n
      c_gpu(i,j) = c_gpu(i,j) + a_gpu(i,k)*b_gpu(k,j)
    enddo
  enddo
enddo
```

LESSON_4 : CUDA C (GPU)

✓ Again: what about performances?

:X

```
// Allocate memory space on the device
REAL *d_a, *d_b, *d_c;
cudaMalloc((void **) &d_a, sizeof(REAL)*nn*nn);
cudaMalloc((void **) &d_b, sizeof(REAL)*nn*nn);
cudaMalloc((void **) &d_c, sizeof(REAL)*nn*nn);

dim3 dimGrid(N_BLOCK, N_BLOCK);
dim3 dimBlock(nn/N_BLOCK, nn/N_BLOCK);

// copy matrix A and B from host to device
cudaMemcpy(d_a, a, sizeof(REAL)*nn*nn, cudaMemcpyHostToDevice);
cudaMemcpy(d_b, b, sizeof(REAL)*nn*nn, cudaMemcpyHostToDevice);

gpu_mm <<< dimGrid, dimBlock >>> (d_a, d_b, d_c, nn);

// Transfer results from device to host
cudaMemcpy(c, d_c, sizeof(REAL)*nn*nn, cudaMemcpyDeviceToHost);
```

LESSON_5 (CPU, GPU)

- ✓ <https://www.youtube.com/watch?v=I40-p9MIG6k&t=3s>
- ✓ Again what about the performance?
 - `nvfortran -Minfo=acc -O2 -stdpar=gpu`

```
do concurrent(j=1:n)
  do k = 1, n
    do i = 1, n
      c(i,j) = c(i,j) + a(i,k)*b(k,j)
    enddo
  enddo
enddo
```

```
do concurrent(j=1:n,i=1:n)
  do k = 1, n
    c(i,j) = c(i,j) + a(i,k)*b(k,j)
  enddo
enddo
```

LESSON_6/1

- ✓ From <https://github.com/ProjectPhysX/OpenCL-Wrapper>
- ✓ <https://www.youtube.com/watch?v=w4HEwdpdTns>

```
int main() {  
    Device device(select_device_with_most_flops()); // compile OpenCL C code for the fastest  
    available device  
  
    const uint N = 1024u; // size of vectors  
    Memory<float> A(device, N); // allocate memory on both host and device  
    Memory<float> B(device, N);  
    Memory<float> C(device, N);  
  
    Kernel add_kernel(device, N, "add_kernel", A, B, C); // kernel that runs on the device  
  
    for(uint n=0u; n<N; n++) {  
        A[n] = 3.0f; // initialize memory  
        B[n] = 2.0f;  
        C[n] = 1.0f;  
    }  
  
    print_info("Value before kernel execution: C[0] = "+to_string(C[0]));  
}
```

LESSON_6/2

```
...  
    A.write_to_device(); // copy data from host memory to device memory  
    B.write_to_device();  
    add_kernel.run(); // run add_kernel on the device  
    C.read_from_device(); // copy data from device memory to host memory  
  
    print_info("Value after kernel execution: C[0] = "+to_string(C[0]));  
    return 0;  
}
```

```
#include "kernel.hpp"  
  
string opcnl_c_container() {  
    kernel void add_kernel(global float* A, global float* B, global float* C) { //  
        equivalent to "for(uint n=0u; n<N; n++) {" , but executed in parallel  
            const uint n = get_global_id(0);  
            C[n] = A[n]+B[n];  
        }  
    }
```


LESSON_7 (GPU)

- ✓ Some intrinsic functions, like `matmul`, has been offloaded to GPU.
 - `nvfortran -Minfo=acc -O2 -acc -gpu=managed -cuda -cudalib`
- ✓ How increases performance with size?

```
#ifdef _OPENACC
    use cutensorex
#endif
.....
    real(my_kind), dimension(:,:), allocatable:: a ! matrix (origin)
    real(my_kind), dimension(:,:), allocatable:: b ! matrix (origin)
    real(my_kind), dimension(:,:), allocatable:: c ! matrix
...
    c = matmul(a,b)
```

LESSON_8 (CPU) /1

✓ MPI program (works only for 2 tasks)

- `mpif90 -O3 mm_mpi.F90`

```
!mpi stuff (setup)
  call MPI_INIT(ierr)
  call MPI_COMM_SIZE(MPI_COMM_WORLD,nprocs,ierr)
  call MPI_COMM_RANK(MPI_COMM_WORLD,myrank,ierr)
```

LESSON_8 (CPU) /2

✓ MPI program (works only for 2 tasks)

- `mpif90 -O3 mm_mpi.F90`

```
!sending a and b elements
  if(myrank.eq.1) then
    call mpi_recv(a(1,1), n*n, MPI_DOUBLE_PRECISION,0,1,      &
                  MPI_COMM_WORLD, status,ierr)
    call mpi_recv(b(1,1), n*n, MPI_DOUBLE_PRECISION,0,2,      &
                  MPI_COMM_WORLD, status,ierr)
  endif
!
  if(myrank.eq.0) then
    call mpi_send(a(1,1), n*n, MPI_DOUBLE_PRECISION,1,1,      &
                  MPI_COMM_WORLD, ierr)
    call mpi_send(b(1,1), n*n, MPI_DOUBLE_PRECISION,1,2,      &
                  MPI_COMM_WORLD, ierr)
  endif
call mpi_barrier(MPI_COMM_WORLD,ierr)
```

LESSON_8 (CPU) /3

✓ MPI program (works only for 2 tasks)

- `mpif90 -O3 mm_mpi.F90`

```
if (myrank == 0) then
  jstart = 1; jend = n/2
endif
!
if (myrank == 1) then
  jstart = n/2+1; jend = n
endif
!
do j=jstart, jend
  do k=1, n
    do i=1, n
      c(i,j) = c(i,j) + a(i,k)*b(k,j)
    enddo
  enddo
enddo
```

LESSON_8 (CPU) /3

✓ MPI program (works only for 2 tasks)

- `mpif90 -O3 mm_mpi.F90`

```
! collecting c elements
if(myrank == 0) then
  call mpi_recv(c(1,n/2+1), n*n/2, MPI_DOUBLE_PRECISION,1,4,&
               MPI_COMM_WORLD, status,ierr)
endif
!
if(myrank == 1) then
  call mpi_send(c(1,n/2+1), n*n/2, MPI_DOUBLE_PRECISION,0,4,&
               MPI_COMM_WORLD, ierr)
endif
call mpi_barrier(MPI_COMM_WORLD,ierr)
...
call MPI_FINALIZE(ierr)
```

LESSON_9

- ✓ Some intrinsic functions, like matmul, has been offloaded to GPU.
 - `mpif90 -O3 -mp mm_mpi.F90`

```
!$OMP PARALLEL DO &
!$OMP DEFAULT(NONE) &
!$OMP PRIVATE(i,j,k) &
!$OMP SHARED(a,b,c,n,jstart,jend)
  do j=jstart, jend
    do k=1, n
      do i=1, n
        c(i,j) = c(i,j) + a(i,k)*b(k,j)
      end do
    end do
  end do
!$OMP END PARALLEL DO
```

Recap/1

✓ Performance in Mflops,

	1024	2048	4096	8192	16384
Lesson_1	-
Lesson_2	
Lesson_3	-
Lesson_4	-
Lesson_5	-
Lesson_7
Lesson_8					
Lesson_9					

Recap/2

- ✓ Size Matters: GPUs and CPUs gives best performance for “relatively” big size.
- ✓ Independently from the paradigm used (OpenACC, OpenMP, CUDA) programmer has to take care and optimize the code (unrolling, blocking)
- ✓ (personal) The best solution doesn't exists. Depends on you experience, knowledge, time devoted to programming, etc. etc.
- ✓ Paradigm can born, change or die in few year. Be ready to change if needed... :-)
- ✓
- ✓ Play a lot, and don't be afraid to make a lot of mistake/errors