

---

# Introduction to High-Performance Computing

Giorgio Amati

Corso di dottorato in Ingegneria Aeronautica e Spaziale 2024

[g.amati@cineca.it](mailto:g.amati@cineca.it) / [g.amaticode@gmail.com](mailto:g.amaticode@gmail.com)

---

---

# CAVEAT

- ✓ This is only an introduction
  - ✓ There's much more under the hood
  - ✓ These 6 lessons can only help you to understand some basic concept
  - ✓ Needs more experience (and so many mistakes) to manage HPC
  - ✓ I'm old guy: almost all examples are in Fortran but the reasons behind performance are (almost) language independent
  - ✓ write to [g.amaticode@gmail.com](mailto:g.amaticode@gmail.com) to have your e-mail
  - ✓ material downloadable from: <https://github.com/gamati01/HPCLessons>
-

---


# Agenda

- ✓ **HPC: What it is?**
  - ✓ **Spoiler...**
  - ✓ Hardware: how it works
  - ✓ Algorithm vs. Implementation
  - ✓ Compiler
  - ✓ Parallel Paradigm
  - ✓ Conclusions & Comments
-

---

## Just for curiosity....

- ✓ Experience of HPC machine?
  - ✓ Fortran, C, C++, everything else?
  - ✓ Parallel paradigm: MPI, OpenMP, OpenACC, OpenMP offload, ....
  - ✓ Linux, Windows, MacOS, (\*NIX)
  - ✓ Are you a Mathematician, a Physicist, a Engineer or a Computer Scientist?
  - ✓ Do you know what is:
    - A Memory System?
    - A Cache?
    - A Floating Point Unit (FPU)?
    - A pipeline?
    - Moore Law?
    - Amdhal Law?
-



---

# HPC: what it is?

From wikipedia:

- ✓ High-performance computing (HPC) uses supercomputers and computer clusters to solve advanced computation problems


Personal definition:

- ✓ It is the overlap of different skills, all devoted to exploit as much as possible the HW performance (both serial and/or parallel, but not limited to supercomputers...)
-


---

# HPC: what it is?

- ✓ These are the main skills for an efficient HPC




# HPC: older than computers?





---

# HPC older than computers



Meteorologist Lewis Fry Richardson, creator of the first dynamic model for weather prediction, proposes the creation of a “forecast factory” that would employ some 64,000 human computers sitting in tiers around the circumference of a giant globe. Each calculator would be responsible for solving differential equations related to the weather in his quadrant of the earth. From a pedestal in the center of the factory, a conductor would orchestrate this symphony of equations by shining a beam of light on areas of the globe where calculation was moving too fast or falling behind.

<https://www.historyofinformation.com/detail.php?id=59>

---

---

# Example: matrix-matrix multiplication

Simple problem: for 2  $n^2$  matrices we have to:

- ✓ compute  $n^3$  products and  $n^3$  sums
- ✓ load  $2 \cdot n^2$  data and to store  $n^3$  data

```
do j = 1, n
  do k = 1, n
    do i = 1, n
      c(i,j) = c(i,j) + a(i,k)*b(k,j)
    enddo
  enddo
enddo
```

- ✓ MM multiplications are used for supercomputing rankings
-

## Example: matrix-matrix multiplication/2

- ✓ Performance can really different, depending on HW, implementation etc....
- ✓ Improvement in performance can be really high....
- ✓ ....or you can easily “depress” performance
- ✓ Performance in Mflops: higher is better

#test	Size	HW	MFlops	Ratio
1	2048	1	201	-
2	2048	1	4870	24x
3	8192	2	361328	1797x
4	8192	2	448923	2233x

## Example: matrix-matrix multiplication/3

Simple problem

- ✓ Performance can really differ, depending on HW, implementation etc....

#test	size	HW	Note	MFlops	Ratio
1	2048	CPU AMD EPYC 7742	cache unfriendly	201	-
2	2048	CPU AMD EPYC 7742	cache friendly	4870	24x
3	8192	GPU A100	do_concurrent	361328	1797x
4	8192	GPU A100	matmul	448923	2233x

---

## Few “facts” about HPC

HPC market is not big enough to survive...


- ✓ SGI
- ✓ Compaq
- ✓ Digital
- ✓ SUN
- ✓ SiCortex
- ✓ MTA
- ✓ CRAY
- ✓ CONVEX
- ✓ CDC
- ✓ Thinking Machine
- ✓ Quadrics/APE

- ✓ IBM
  - Power3/4/.../9
- ✓ Intel
  - Itanium
  - Phi
- ✓ HP
- ✓ NVIDIA
- ✓ FUJITSU
- ✓ AMD
  - Opteron
- ✓ NEC
  - SX6


# Few “facts” about HPC

Top500 list: [June 2003](#) vs [November 2022](#)

Processor Generation System Share



Processor Generation System Share




# Moore's Law

In his article in 1965, Moore (Intel co-founder) planned the increase of the # of transistors up to 1975.


***“With unit cost falling as the number of components per circuit rises, by 1975 economics may dictate squeezing as many as 65,000 components on a single silicon chip”***

He stated a 2x increment of transistors every 18 Month.

- ✓ This law is still valid now (in some form): to “survive” in the market HW firms must follow this law
- ✓ Now “transistor shrinking” is much harder: we are near to quantum effects
- ✓ New “ideas” must be found




# Follow the money! (from D.Reed@ISC2023)





# Follow the money! (from D.Reed@ISC2023)



# Follow the money! (from D.Reed@ISC2023)



AI, cloud, and silicon innovation  
Follow the money and the culture ...

Hardware unicorns and AI startups

- Cerebras, GraphCore, Groq, Hailo
- SambaNova, Wave Computing, ...



Google TPU4 (operational 2020)

- 4096 units per "pod" (1.1 exaops)
- 3-D twisted torus **optical interconnect**


AWS Graviton3

- 64 ARM Neoverse V1 cores, 7 chiplet design
- 55 billion transistors, DDR5 memory, PCIe5

Microsoft Azure (XCG legacy, in part)

- Ampere ARM and Project Catapult/Brainwave
- \$10B+ OpenAI investment

Ampere One (192 cores, Bfloat16)



Microsoft



IMAGINE

TOMORROW

Old stuff....



2023/1/20 09:13




2023/1/20 09:13



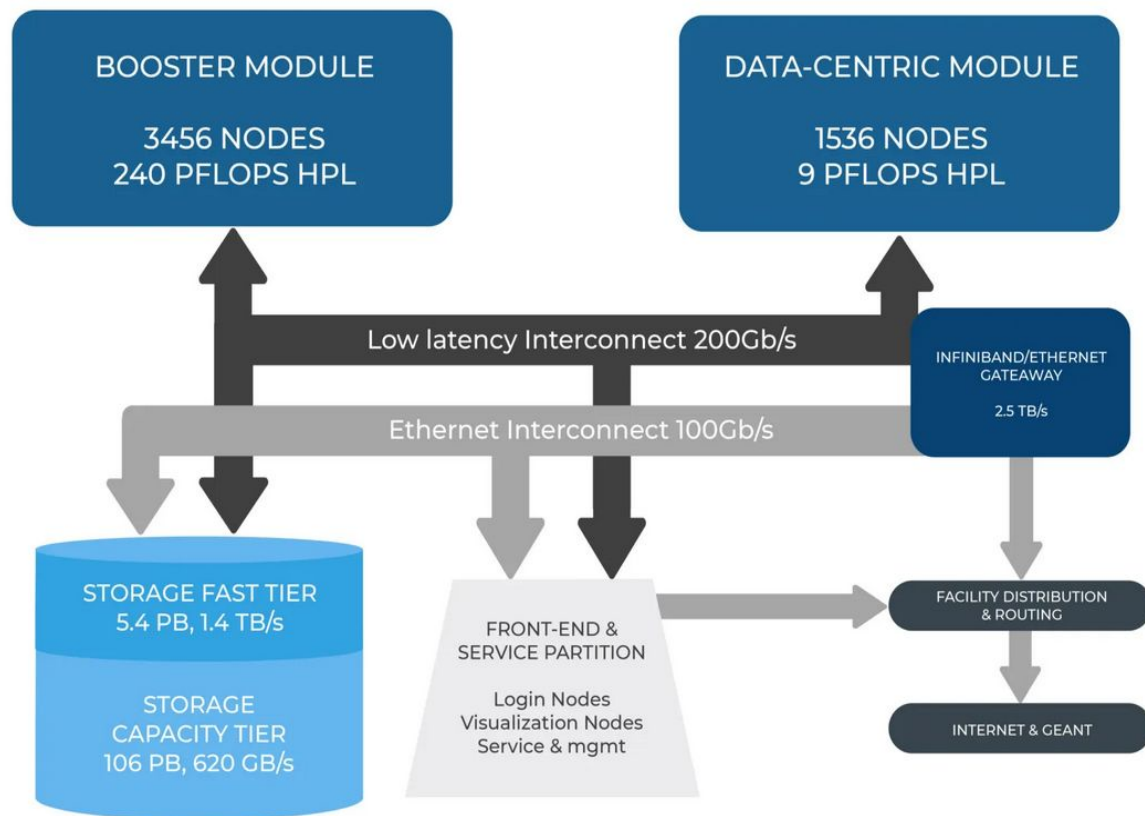
2023/1/20 09:13






# Leonardo: main figures

- ✓ 1536 CPU-based nodes
  - 172032 cores
- ✓ 3456 GPU-based nodes
  - 13824 GPU
  - 110592 cores
- ✓ 155 Racks
  - 16 CPU racks
  - 116 GPU racks
  - 12 I/O racks
  - 1 System racks
  - **About 300'000 Kg!**
- ✓ Power Requirements
  - HPL: ~ 8.0 MW
  - Operational: ~ 6.0 MW



# CPU vs. GPU



- ✓ Optimized for low latencies
- ✓ Huge caches
- ✓ Control logic for out-of-order and speculative execution
- ✓ **Targets on general-purpose applications**

- ✓ Optimized for data parallel throughput
- ✓ Memory latency tolerant
- ✓ More transistors dedicated to computations
- ✓ **Targets on special applications**

---

# Why GPUs?

## ✓ Pro

- GPU more powerful: 1 GPU ~ 10x CPU (Peak Mflops)
- GPU ask for less space: for same performance CPU ask for ~3x more racks
- GPU are less expensive: for same peak performance CPU are ~2x more expensive
- GPU asks for (relative) less power: for the same peak performance CPU needs ~4x more energy

## ✓ Cons

- GPU are less flexible respect CPU
  - Some algorithm are not GPU-friendly
  - There's no a common programming model between different vendors
  - Porting to GPU is expensive and error-prone procedure
-

---

## Recap

- ✓ Different skills are required to achieve “good” performance
  - ✓ Performance is not only a problem of the right (powerful) HW
  - ✓ HW evolution is driven by mass market
  - ✓ All firms devoted only to HPC have not survived to the market
  - ✓ User should (or must) be flexible enough to follow HW & SW evolution
  - ✓ A Correct code could be efficient or not. With different order of magnitude!
  - ✓ Today any processor is a parallel one
    - To have a parallel code doesn't mean to have an efficient one
  - ✓ To be fast is secondary respect to be correct
    - “Premature optimization is the root of all evils” (D. Knuth)
  - ✓ But you'll must face soon optimization issues
    - 1 way to go fast, 100 ways to go slow!
  - ✓ Today CPU/GPUs can have order of 100'000'000'000 transistors
-