

---

# Introduction to High-Performance Computing

Giorgio Amati

Corso di dottorato in Ingegneria Aeronautica e Spaziale 2024

[g.amati@cineca.it](mailto:g.amati@cineca.it) / [g.amaticode@gmail.com](mailto:g.amaticode@gmail.com)

---

---

# Agenda

- ✓ **HPC: What it is?**
  - ✓ **Spoiler...**
  - ✓ **Hardware: how it works**
  - ✓ Algorithm vs. Implementation
  - ✓ Compiler
  - ✓ Parallel Paradigm
  - ✓ Conclusions & Comments
-

---

# HPC: what it is?

- ✓ These are the main skills for an efficient HPC

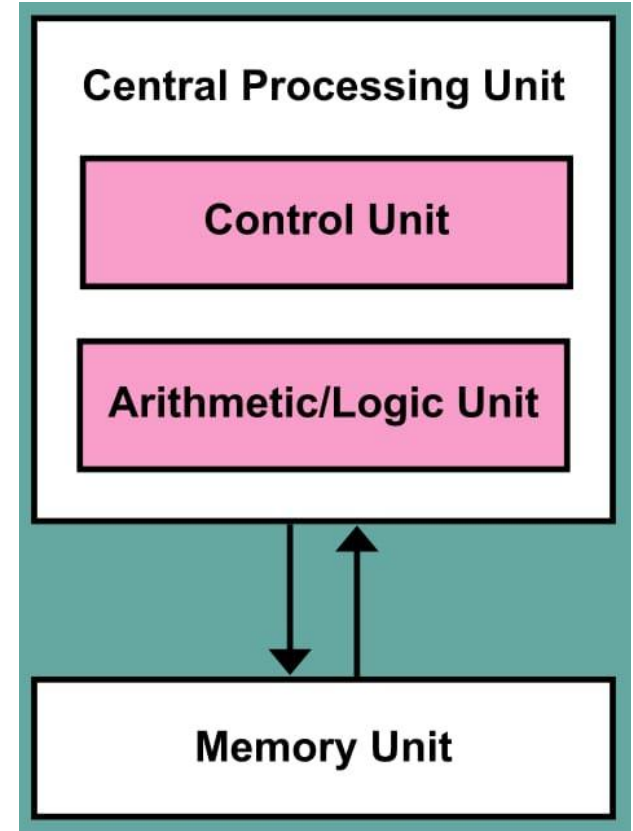


# Von Neumann's model

- ✓ Independent Systems that talk each other exchanging data
- ✓ Data and instruction are stored in the same area
- ✓ Different, and independent, units for instruction (Control Units) and data (ALU units)


## Two key points

- ✓ Data to move back and forth from main memory
- ✓ Instruction (e.g. Flops) to perform moved from Memory



# BW & Flops: evolution

✓ Hic sunt leones! (From Hennessy & Patterson, 2011)



---

## Two metrics...

### Latency

- ✓ Time need to complete one operation (once all the input data are available)
- ✓ Lower is better

### Throughput

- ✓ How many operations can complete according to a time interval
- ✓ Higher is better

### Example:

- ✓ Elevator: low latency, low throughput
  - ✓ Escalator: high latency, high throughput
-

---

# Floating point units

The big increase in Performance (or Flops) was due to:

- ✓ optimization of the Floating Point Unit (FPU)
- ✓ a big number of FPU/Cores in a CPU (thanks to Moore's Law)


A generic operation can be splitted in different (independent) stages: e.g (but its a simplification):

- ✓ Fetch
- ✓ Decode
- ✓ Execute

- ✓ Let's assume that each stage can be completed in a single clock cycle
  - ✓ Let's assume that the operation is a floating point operation (e.g. a sum or a product)
-

# Floating point units

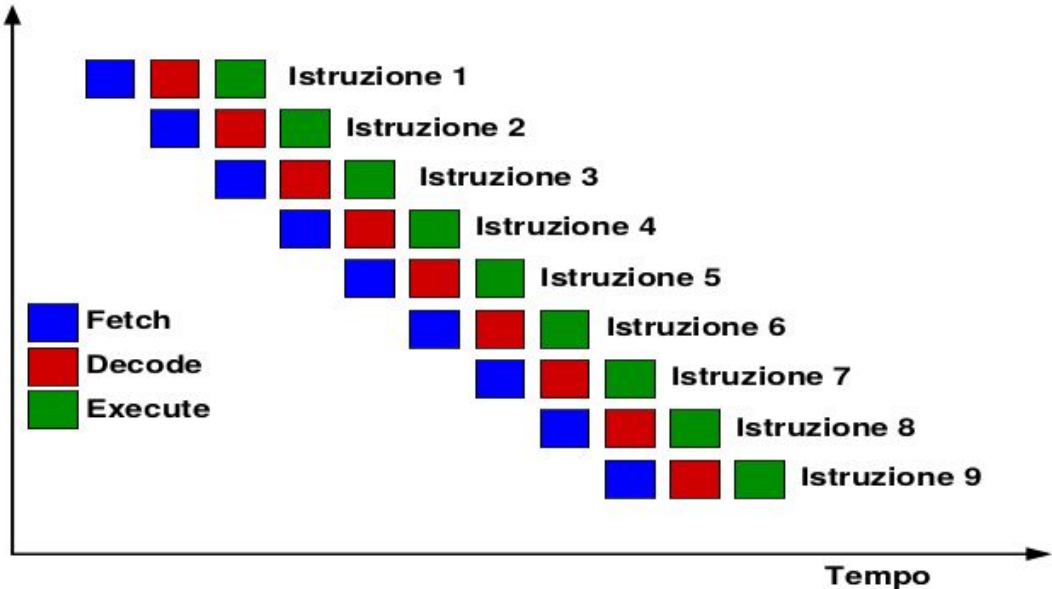
- ✓ Simple FPU: an instruction completed every 3 cycles
- ✓ 1 Flop every 3 cycle: we need to have 2 load+1store every 3 cycles
- ✓ Latency = 3 cycles






# Pipelining

- ✓ Pipelined FPU: all the three stages are independent (different transistors)
- ✓ 1 Flop computed every cycle: we need  $2load+1$  store every cycle
- ✓ at least 3 independent instructions (intr #1, #2, #3)
- ✓ Latency=3 cycles




# Superpipelining

- ✓ Superpipelined FPU: Further decomposition in more stages
- ✓ We can (ideally) increase by a factor 2 the frequency
- ✓ 1 Flop computed every cycle: we need 2 load+1 store every cycle (but doubled frequency)
- ✓ Latency = 6 cycles




# Superscalar

- ✓ More transistor → More FPU
- ✓ 2 Flop computed every cycle: we need 4 load+2 store every cycle
- ✓ 2 independent instructions **every cycle**
- ✓ Latency = 6 cycles



# Out Of Order Execution

- ✓ “On the fly” reordered instruction in order to avoid “bubble” in the pipeline
- ✓ Same request of Superscalar but better performance



# Pipeline Recap

With pipelining we see a big increase in performance (e.g. #Flops) but we need


- ✓ for Superscalar FPU
  - 12 times more independent instruction respect a non-pipelined FPU
  - 12 times more data to perform all requested operations
- ✓ Both Mem subsystem & Implementation can help

	FPU	Pipelined	Superpipelined	Superscalar	OOO
stages	1	3	6	6	6
frequency	1	1	1/2	1/2	1/2
performance	1	3	6	12	12
Requested BW	1/3	1	2	4	4

# FPU: Vector Unit

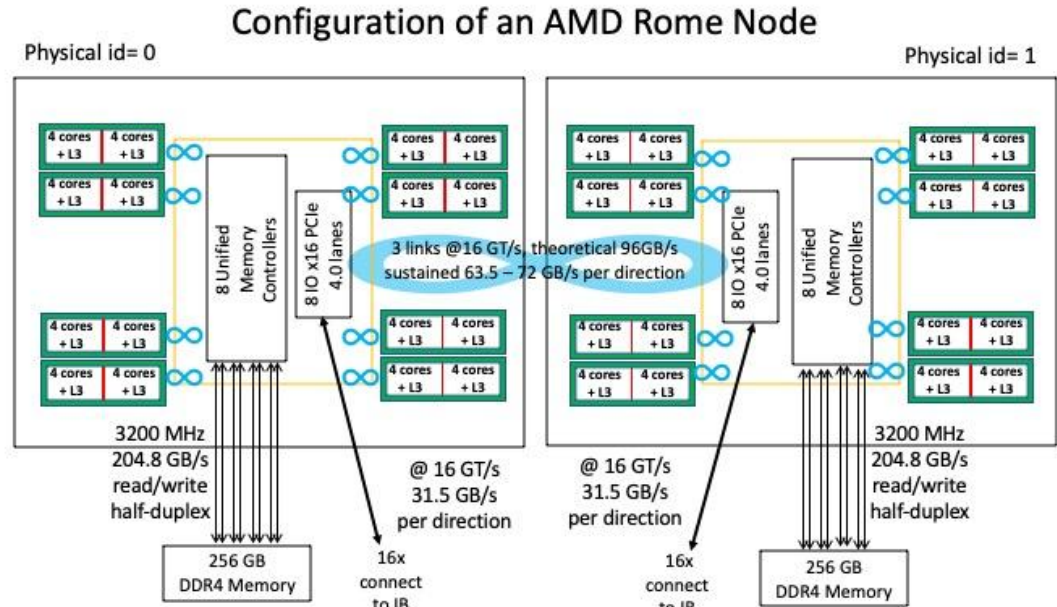
Further improvement: operation (e.g. sum) are performed using a vector of 512 byte (4 Double precision or 8 Single precision, AVX-512)

- ✓ 8x improvement respect single scalar unit
- ✓ 8x request of independent instruction
- ✓ x8x more BW requirement




# A today CPU

- ✓ [https://www.nas.nasa.gov/hecc/support/kb/amd-rome-processors\\_658.html](https://www.nas.nasa.gov/hecc/support/kb/amd-rome-processors_658.html)
- ✓ Order of 40 Billion Transistors
- ✓ 64 Cores: each core
  - 2 FMA units at 256-bit
  - 4 x86 instruction per cycle
  - 4 flops per cycle



# Moore's Law




Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten  
New plot and data collected for 2010-2015 by K. Rupp



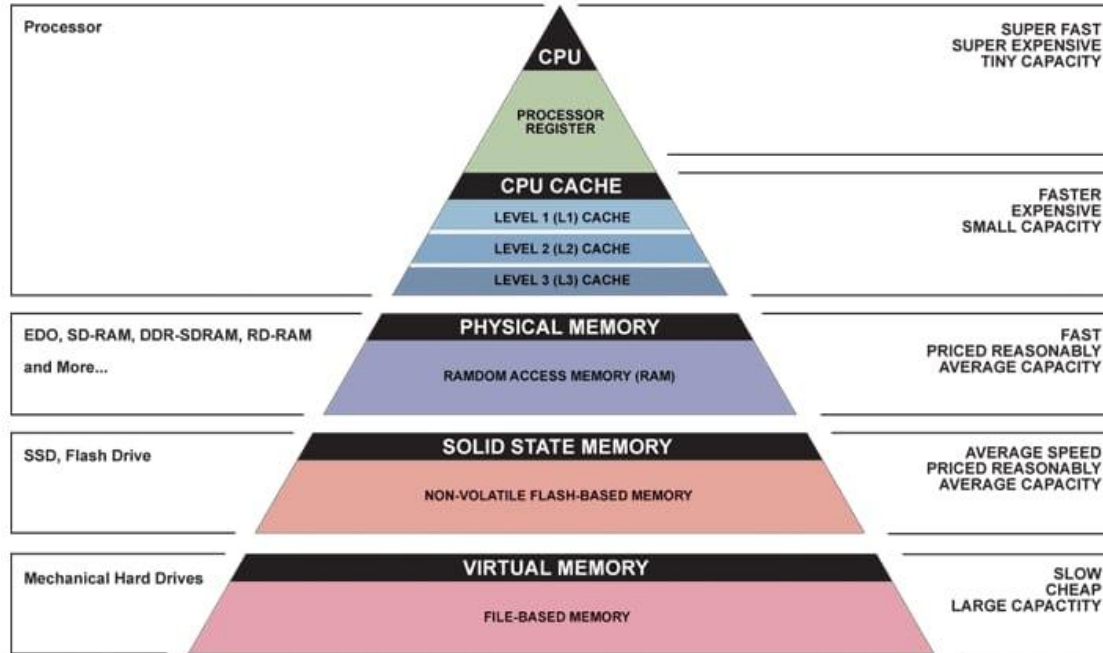
# BW & Flops: evolution

✓ Hic sunt leones! (From Hennessy & Patterson, 2011)



# Memory Subsystem

Memory subsystem: hierarchical structure



▲ Simplified Computer Memory Hierarchy  
Illustration: Ryan J. Leng

---

# Memory Subsystem: velocity vs Size

[https://www.nas.nasa.gov/hecc/support/kb/amd-rome-processors\\_658.html](https://www.nas.nasa.gov/hecc/support/kb/amd-rome-processors_658.html)

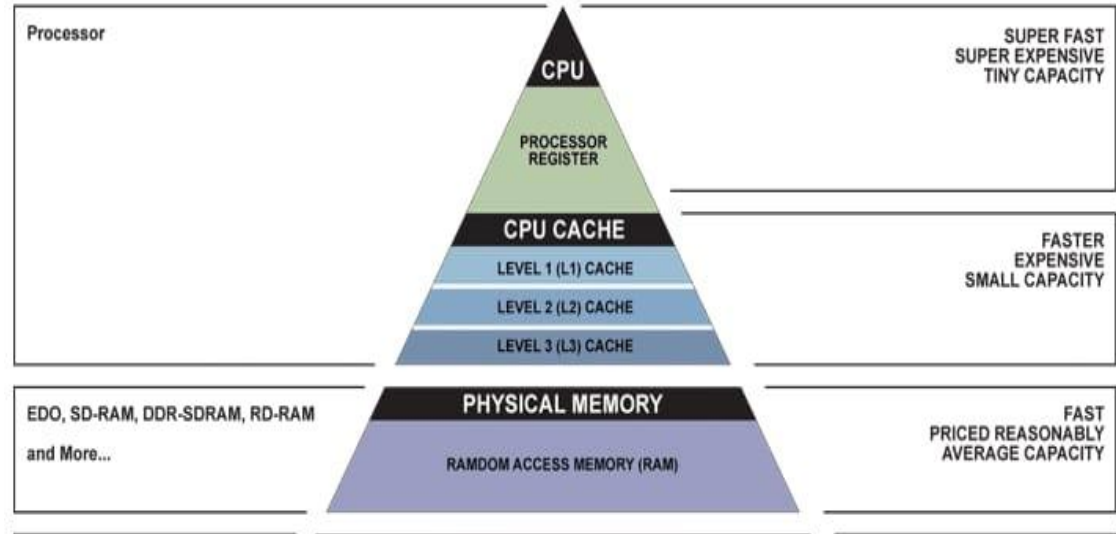

Unit	Size	Latency
Register	~32/128	~1 Cycles
Cache Level 1 (Data)	32 KB	~7 Cycles
Cache Level 1 (istr.)	32 KB	~7 Cycles
Cache Level 2	512 KB (per Core)	~12 Cycles
Cache Level 3	16MB (per 4Core)	~36 Cycles
HMB	~64 GB	~100/1000 Cycles
RAM	> 100 GB	~ 1000/10000 Cycles

---

# Memory Subsystem: velocity

Data movement is expensive: RAM latency is order of 1'000 of cycles

- ✓ Intermediate (cache) levels to hide latency
- ✓ All data are stored in RAM memory. In cache we store intermediate results or copies



---

# Principle of Locality

- ✓ Locality is the key point of the memory subsystem structure.
- ✓ Locality is crucial to achieve performance. Without exploiting locality performance are very very low

**locality is the tendency of a processor to access the same set of memory locations repetitively over a short period of time**

- ✓ Two different localities are implemented:

**Spatial (or data) locality**

**Temporal locality**

[https://en.wikipedia.org/wiki/Locality\\_of\\_reference](https://en.wikipedia.org/wiki/Locality_of_reference)

---

---

# Space Locality

Usually a program use data that are “near” each other.

- ✓ If at  $t$ , I need data in location  $A$
- ✓ at time  $t+1$ , I'll look for a data in location  $B$  “near” to  $A$

Solution

- ✓ When I've to “access” to location  $A$ , I can try to “access” also to location  $B$ , to hide data access latency (prefetch)
-

---

# Temporal locality

Usually a program re-use data in a short time-frame

- ✓ If at time  $t$ , I need data in location  $A$
- ✓ at time  $t+1$ ,  $t+2$ ,  $t+3$ ,.... I'll need again the data stored in location  $A$


Solution:

- ✓ Store the data in location  $A$  “somewhere” before write it back to RAM
-

# Exploiting locality: Cache line

Data are moved from RAM memory to cache in a cache-line: usually 64 byte wide.

- ✓ RAM latency is order of 1'000 Cycles
  - ✓ Cache latency is order of 10 Cycles
- 
- ✓ To exploit Spatial locality
  - ✓ FPU “see” the Cache latency not the RAM one





---

# Temporal locality

Data used many time are stored in cache instead of being copied (every time) back in RAM.

This applies in particular to

- ✓ Intermediate Results
- ✓ Constant

**Remember:** data in cache are copies of the “original” data in RAM, the OS is “obliged” to synchronize between cache and RAM

---

---

# Cache friendly vs. cache unfriendly

## Cache unfriendly

1. Look for A(1)
2. cache miss (compulsory)
3. load from RAM to cache
4. move data to register
5. perform the operation
6. write back the result
7. look for A(2)
8. cache miss
9. load from RAM to cache
10. move data to register
11. perform the operation
12. write back the result
13. look for A(3)
14. cache miss
15. ....

## Cache friendly

1. Look for A(1)
  2. cache miss (compulsory)
  3. load from RAM to cache
  4. move data to register
  5. perform the operation
  6. write back the result
  7. look for A(2)
  8. cache hit
  9. move data to register
  10. perform the operation
  11. write back the result
  12. look for A(3)
  13. cache hit
  14. move data to register
  15. ....
-

---

# Cache friendly vs. cache unfriendly

## Cache unfriendly


1. Look for A(1)
2. **cache miss**
3. **load from RAM to cache**
4. move data to register
5. perform the operation
6. write back the result
7. look for A(2)
8. **cache miss**
9. **load from RAM to cache**
10. move data to register
11. perform the operation
12. write back the result
13. look for A(3)
14. **cache miss**
15. ....

## Cache friendly

1. Look for A(1)
  2. **cache miss**
  3. **load from RAM to cache**
  4. move data to register
  5. perform the operation
  6. write back the result
  7. look for A(2)
  8. **cache hit**
  9. move data to register
  10. perform the operation
  11. write back the result
  12. look for A(3)
  13. **cache hit**
  14. move data to register
  15. ....
-


## Data allocation: C

- ✓ Programmer style/languages could have an impact in the performance
- ✓ Memory is a linear Array: how map matrix  $A[i][j]$ ?
- ✓ C is row oriented
- ✓ In C a matrix is a vector of vectors
- ✓  $A[1][1], A[1][2] \rightarrow \text{stride}=1$
- ✓  $A[1][1], A[2][1] \rightarrow \text{stride}=n$




## Data allocation: Fortran

- ✓ Programmer style/languages could have an impact in the performance
- ✓ Memory is a linear Array: how map matrix  $A(i, j)$ ?
- ✓ Fortran is column oriented
- ✓  $A(1,1), A(1,2) \rightarrow \text{stride}=n$
- ✓  $A(1,1), A(2,1) \rightarrow \text{stride}=1$



## Data allocation: effects

The matrix (left) is allocated in a completely different way between Fortran and C



---

# matrix-matrix multiplication (Fortran)

Both code are correct : but which has unitary stride?

```
do j = 1, n
  do k = 1, n
    do i = 1, n
      c(i,j) = c(i,j)+a(i,k)*b(k,j)
    enddo
  enddo
enddo
```

```
do i = 1, n
  do k = 1, n
    do j = 1, n
      c(i,j) = c(i,j)+a(i,k)*b(k,j)
    enddo
  enddo
enddo
```

---

# matrix-matrix multiplication (C)

Both code are correct : but which has unitary stride?

```
for (i = 0; i < nn; i++)  
    for (k = 0; k < nn; k++)  
        for (j = 0; j < nn; j++)  
            c[i][j]=c[i][j]+a[i][k]*b[k][j];
```

```
for (j = 0; j < nn; j++)  
    for (k = 0; k < nn; k++)  
        for (i = 0; i < nn; i++)  
            c[i][j]=c[i][j]+a[i][k]*b[k][j];
```



## Some figure

Performance in Mflops (higher is better)

- ✓ HW: Intel(R) Xeon(R) Platinum 8260 CPU @ 2.40GHz
- ✓ Compiler intel 2021.5
  - `ifort -O1`
  - `icc -O1`
- ✓ Matrix size: 2048 (96 MB)

Index order	Language	Mflops
i,k,j	Fortran	2870
j,k,i	Fortran	129
i,k,j	C	181
j,k,i	C	2373

## Example: matrix-matrix multiplication

- ✓ Performance can really different, depending on HW, implementation etc....
- ✓ Improvement in performance can be really high....
- ✓ ....or you can easily “depress” performance
- ✓ Performance in Mflops: higher is better

#test	Size	HW	MFlops	Ratio
1	2048	1	201	-
2	2048	1	4870	24x
3	8192	2	361328	1797x
4	8192	2	448923	2233x

# Size matters!!!!

- ✓ Performance, in MFlops, for different problem size: from  $n=100^2$  to  $n=4000^2$ 
  - using cache friendly access (stride=1)
  - using cache un-friendly access (stride=n)

	256	512	769	1024	1536	2048
stride=n	268	367	335	113	71	49
stride=1	1073	2643	2828	2290	2273	1683
Ratio	4.0x	7.2x	8.4x	20x	32x	34x

---

## Data allocation: SoA or AoS?

The way we allocate data force the way we access to them!

e.g.: Velocity for a cartesian grid

- ✓ Structure of Arrays (SoA):
- ✓ Arrays of Structures (AoS):

```
! vel=1,2,3 where 1=vx, 2=vy, 3=vz)
a(i,j,k,vel) = ...
```

```
! vel=1,2,3 where 1=vx, 2=vy, 3=vz)
a(vel,i,j,k) = ...
```

- ✓ Which is faster?

# Exploiting Cache usage: blocking

## Matrix Transpose

- ✓ Or load or store with not unitary stride
- ✓ Split in different blocks, little enough to fit in cache

```
do j = 1, n
  do i = 1, n
    B(i,j) = A(j,i)
  enddo
enddo
```

```
do jj = 1, n, step
  do ii = 1, n, step
    do j = jj, jj+step-1
      do i = ii, ii+step-1
        B(i,j) = A(j,i)
      enddo
    enddo
  enddo
enddo
```

# Exploiting performance: unrolling

Unrolling exploits instruction independence

- ✓ Allows to increase the number of data streams
- ✓ be careful to index extremes
- ✓ can introduce systematic cache trashing

```
do j = 1, n
  do k = 1, n
    do i = 1, n
      c(i,j) = c(i,j)+a(i,k)*b(k,j)
    enddo
  enddo
enddo
```

```
do j = 1, n, 4
  do k = 1, n
    do i = 1, n
      c(i,j+0)=c(i,j+0)+a(i,k)*b(k,j+0)
      c(i,j+1)=c(i,j+1)+a(i,k)*b(k,j+1)
      c(i,j+2)=c(i,j+2)+a(i,k)*b(k,j+2)
      c(i,j+3)=c(i,j+3)+a(i,k)*b(k,j+3)
    enddo
  enddo
enddo
```

# Exploiting performance: unrolling

Unrolling exploits instruction independence

- ✓ Allows to increase the number of data streams
- ✓ be careful to index
- ✓ can introduce systematic cache trashing

```
do j = 1, n
  do k = 1, n, 2
    do i = 1, n
      c(i,j)=c(i,j)          &
      +a(i,k+0)*b(k+0,j) &
      +a(i,k+1)*b(k+1,j) &
    enddo
  enddo
enddo
```

```
do j = 1, n
  do k = 1, n
    do i = 1, n, 4
      c(i+0,j)=c(i+0,j)+a(i+0,k)*b(k,j)
      c(i+1,j)=c(i+1,j)+a(i+1,k)*b(k,j)
      c(i+2,j)=c(i+2,j)+a(i+2,k)*b(k,j)
      c(i+3,j)=c(i+3,j)+a(i+3,k)*b(k,j)
    enddo
  enddo
enddo
```

## Some further figure

Performance in Mflops (higher is better)

- ✓ HW: Intel(R) Xeon(R) Platinum 8260 CPU @ 2.40GHz
- ✓ Compiler intel 2021.5
  - `ifort -O1`
- ✓ Matrix size: 2048 (96 MB)

Optimization	Mflops
<b>simple</b>	<b>2210</b>
<b>blocking (32)</b>	<b>2928</b>
<b>unrolling x 4 (j)</b>	<b>3267</b>
<b>unrolling x 4 (i)</b>	<b>2530</b>
<b>unrolling x 2 (k)</b>	<b>3420</b>
<b>All together</b>	<b>5653</b>



# Exploiting performance: loop merging

```
do 1000 z=1,nz
  k3=beta(z)
  do 1000 y=1,ny
    k2=eta(y)
    do 1000 x=1,nx/2
      hr(x,y,z,1)=hr(x,y,z,1)*norm
      hi(x,y,z,1)=hi(x,y,z,1)*norm
      hr(x,y,z,2)=hr(x,y,z,2)*norm
      hi(x,y,z,2)=hi(x,y,z,2)*norm
      hr(x,y,z,3)=hr(x,y,z,3)*norm
      hi(x,y,z,3)=hi(x,y,z,3)*norm
    1000 continue
  c
  do 2000 z=1,nz
    k3=beta(z)
    do 2000 y=1,ny
      k2=eta(y)
      do 2000 x=1,nx/2
        ur(x,y,z,1)=ur(x,y,z,1)*norm
        ur(x,y,z,1)=ur(x,y,z,1)*norm
        ur(x,y,z,2)=ur(x,y,z,2)*norm
        ui(x,y,z,2)=ui(x,y,z,2)*norm
        ui(x,y,z,3)=ui(x,y,z,3)*norm
        ui(x,y,z,3)=ui(x,y,z,3)*norm
      2000 continue
    c
    do 3000 z=1,nz
      k3=beta(z)
      do 3000 y=1,ny
        k2=eta(y)
        do 3000 x=1,nx/2
          k1=alfa(x,1)
          k_quad=k1*k1+k2*k2+k3*k3+k_quad_cfr
          k_quad=1./k_quad
          sr=k1*hr(x,y,z,1)+k2*hr(x,y,z,2)+k3*hr(x,y,z,3)
          si=k1*hi(x,y,z,1)+k2*hi(x,y,z,2)+k3*hi(x,y,z,3)
          hr(x,y,z,1)=hr(x,y,z,1)-sr*k1*k_quad
          hr(x,y,z,2)=hr(x,y,z,2)-sr*k2*k_quad
          hr(x,y,z,3)=hr(x,y,z,3)-sr*k3*k_quad
          hi(x,y,z,1)=hi(x,y,z,1)-si*k1*k_quad
          hi(x,y,z,2)=hi(x,y,z,2)-si*k2*k_quad
          hi(x,y,z,3)=hi(x,y,z,3)-si*k3*k_quad
          k_quad_cfr=0.
        3000 continue
      c
    1000 continue
  c
```

!! primo loop

!! secondo loop

!! terzo loop

FPU need a lot of data to be  
“full”.

Sometimes a single loop has  
not enough instruction to be  
efficient

It increase the number on  
independent instructions

```
do 1000 z=1,nz
  k3=beta(z)
  do 1000 y=1,ny
    k2=eta(y)
    do 1000 x=1,nx/2
      hr(x,y,z,1)=hr(x,y,z,1)*norm
      hi(x,y,z,1)=hi(x,y,z,1)*norm
      hr(x,y,z,2)=hr(x,y,z,2)*norm
      hi(x,y,z,2)=hi(x,y,z,2)*norm
      hr(x,y,z,3)=hr(x,y,z,3)*norm
      hi(x,y,z,3)=hi(x,y,z,3)*norm
    1000 continue
  c
  do 2000 z=1,nz
    k3=beta(z)
    do 2000 y=1,ny
      k2=eta(y)
      do 2000 x=1,nx/2
        ur(x,y,z,1)=ur(x,y,z,1)*norm
        ur(x,y,z,1)=ur(x,y,z,1)*norm
        ur(x,y,z,2)=ur(x,y,z,2)*norm
        ui(x,y,z,2)=ui(x,y,z,2)*norm
        ui(x,y,z,3)=ui(x,y,z,3)*norm
        ui(x,y,z,3)=ui(x,y,z,3)*norm
      2000 continue
    c
    do 3000 z=1,nz
      k3=beta(z)
      do 3000 y=1,ny
        k2=eta(y)
        do 3000 x=1,nx/2
          k1=alfa(x,1)
          k_quad=k1*k1+k2*k2+k3*k3+k_quad_cfr
          k_quad=1./k_quad
          sr=k1*hr(x,y,z,1)+k2*hr(x,y,z,2)+k3*hr(x,y,z,3)
          si=k1*hi(x,y,z,1)+k2*hi(x,y,z,2)+k3*hi(x,y,z,3)
          hr(x,y,z,1)=hr(x,y,z,1)-sr*k1*k_quad
          hr(x,y,z,2)=hr(x,y,z,2)-sr*k2*k_quad
          hr(x,y,z,3)=hr(x,y,z,3)-sr*k3*k_quad
          hi(x,y,z,1)=hi(x,y,z,1)-si*k1*k_quad
          hi(x,y,z,2)=hi(x,y,z,2)-si*k2*k_quad
          hi(x,y,z,3)=hi(x,y,z,3)-si*k3*k_quad
          k_quad_cfr=0.
        3000 continue
      c
    1000 continue
  c
```

!! primo loop

!! secondo loop

!! terzo loop



---

# Exploiting performance: loop splitting

The opposite of loop merging: sometimes a loop is so “fat” that the compiler is not able to efficiently fetch enough independent instruction: It work on a finite window of lines and cannot have a complete vision of the operation to be performed

- ✓ Vectorization
  - ✓ It's a “try and error” procedure
-

# Spoiler: Distributed & Shared Memory systems



---

## (some) Keywords

- ✓ **Bandwidth:** the rate of data transfer from/to CPU and Memory Subsystem
  - ✓ **FMA FPU:** Fused Multiply and ADD FPU
  - ✓ **Latency:** time delay between the start and the end of the operation (usually in cycles)
  - ✓ **Cache:** is a hardware or software component that stores data so that future requests for that data can be served faster
  - ✓ **Cache hit/miss:** when a data is found/not found in the cache
  - ✓ **Cache Trashing:** systematic trashing of a cache line due to different data stream insisting on the same cache line
  - ✓ **Stride:** distance in memory between two different consecutive access
  - ✓ **Spatial locality (also data locality):** data/instruction elements within relatively close storage locations
  - ✓ **Temporal locality:** data/instruction elements that will be used again soon
-

---

# RECAP

- ✓ Today HW is very performing if, and only if, you perform
    - an efficient use of the FPU's
    - an efficient use of the Memory sub-system
  
  - ✓ From user point of view
    - Take care of data allocation
    - Access at unitary stride if possible
    - Exploit independent instructions
    - Exploit data streams
    - **A program can be written in many different way, all correct but with very different performance figures!**
-

---

## Topic not covered here

**Virtual memory**, or **virtual storage**, is a memory management technique that provides an "idealized abstraction of the storage resources that are actually available on a given machine which "creates the illusion to users of a very large (main) memory". The computer's operating system, using a combination of hardware and software, maps memory addresses used by a program, called *virtual addresses*, into *physical addresses* in computer memory. Main storage, as seen by a process or task, appears as a contiguous address space or collection of contiguous segments. The operating system manages virtual address spaces and the assignment of real memory to virtual memory. Address translation hardware in the CPU, often referred to as a memory management unit (MMU), automatically translates virtual addresses to physical addresses. Software within the operating system may extend these capabilities, utilizing, e.g., disk storage, to provide a virtual address space that can exceed the capacity of real memory and thus reference more memory than is physically present in the computer....

[https://en.wikipedia.org/wiki/Virtual\\_memory](https://en.wikipedia.org/wiki/Virtual_memory)

---

---

## Some Books/Reference

- ✓ Charles Severance; Kevin Dowd “High Performance Computing”, O’Reilly, ISBN 13:9781565923126
- ✓ John L. Hennessy, David A. Patterson , “Computer Architecture: A Quantitative Approach” Morgan Kaufmann; ISBN-10 : 0128119055
- ✓ John L. Hennessy, David A. Patterson , “Computer Organization and Design: The Hardware/Software Interface”, Morgan Kaufmann;
- ✓ D. Goldberg, [“What Every Computer Scientist Should Know About Floating-Point Arithmetic”](#)
- ✓ U. Drepper: [“What Every Programmer Should Know About Memory”](#)