Introduction to
High-Performance Computing

Giorgio Amati

Corso di dottorato in Ingegneria Aeronautica e Spaziale 2024

g.amati@cineca.it / g.amaticode@gmail.com

Agenda

✓ **HPC: What it is?**
✓ **Spoiler…**
✓ **Hardware: how it works**
✓ **Algorithm vs. Implementation**
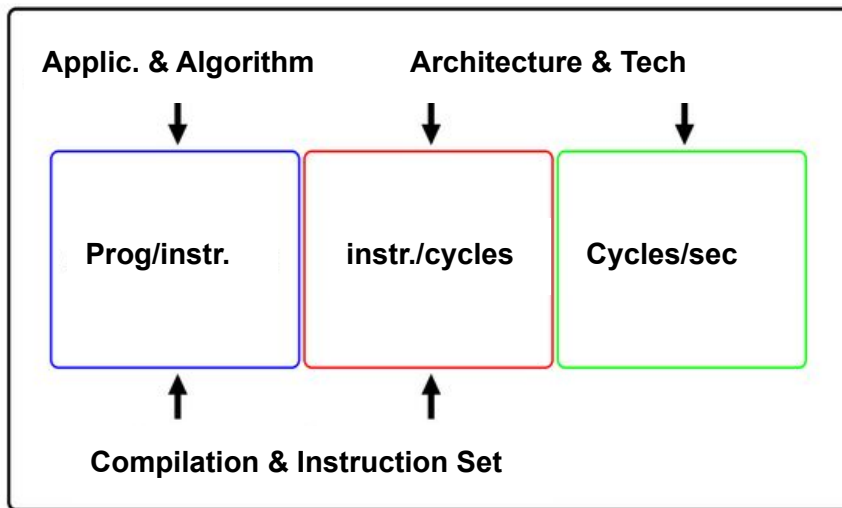✓ Compiler
✓ Parallel Paradigm
✓ Conclusions & Comments

✓ These are the main skills for an efficient HPC

Who is in charge of what?

High performance is the joint effort of different "players"

✓ **Programmer**: in charge of the choice of the algorithm
✓ **Compiler**: in charge of the traduction in instructions
✓ **HW**: in charge of execute the instructions

Algorithm

Definition:

a process or set of rules to be followed in calculations or other problem-solving operations, especially by a computer.....

✓ From previous lesson we have learned that, for a computer, different instruction can have different impact in time
✓ We are not talking about "correctness" (when time is not an issue) but about "computational complexity"
✓ As example
   ■ $O(N^2) << O(N^4)$ for N big enough (still to define)
✓ The choice, at the beginning, of an inefficient algorithm can be really "dangerous"!

Complexity

In computer science, the computational complexity or simply complexity of an algorithm is the amount of resources required to run it. Particular focus is given to computation time (generally measured by the number of needed elementary operations) and memory storage requirements.

- ✓ Each algorithm has it own complexity, i.e. the number of instruction to be performed to complete the work
- ✓ If **N** is the size of the problem **$O(N^?)$** is the problem complexity
    - ■ MMM → $O(N^3)$
    - ■ MMM (Stressen) → $O(N^{2.80})$
- ✓ Hint:
    - ■ Try to understand complexity of your algorithm
    - ■ Verify if your assumption is correct: it scale as it is supposed to?
    - ■ Stress to algorithm for "real-size" problem: usually development is done using very little problem (in size). to little to be "killed" by high complexity (e.g. cache effects)
    - ■ For little problems size the prefactor could "hide" real performance

Hardware evolution (from 1947 to 1985)

✓ From: More Programming Pearls: Confessions of a Coder, Bentley

| HW	Year	Mflops
Manchester Mark 1	1947	0.0002
IBM 701	1954	0.003
IBM Stretch	1960	0.3
CDC 6600	1964	2
CDC 7600	1969	5
Cray-1	1976	50
Cray-2	1985	125

Algorithm Evolution

Algorithm evolution (from 1945 to 1978) for 3D elliptic equation (e.g. pressure)

✓ From: More Programming Pearls: Confessions of a Coder, Bentley

Algorithm	Year	Complexity
Gaussian Elimination	1947	N^7
SOR (suboptimal)	1954	8N^5
SOR (optimal)	1960	8N^4*log(N)
Cyclic Reduction	1964	8N^3*log(N)
Multigrid	1969	60N^3

Which is more important?

Which the limit?

In HPC is mandatory to know which is the limit do the optimization process?

- ✓ How far I am for the limit?
- ✓ I'm inside my time constraint?

A HW system has 2 key figures

- ✓ How many floating point operation I can perform? → FLOPS
- ✓ How many data i can move back and forth → BW

Arithmetic Intensity of a code

- ✓ A.I = ratio #Flops/#Byte moved

Roofline Model

✓ https://www.nersc.gov/assets/Uploads/Tutorial-ISC2019-Intro-v2.pdf

Roofline Model

How increase performance?

✓ Reduce data movement
✓ Increase Floating point operations (?)

Performance evolution/Intel

CPU (codename)	Clock Frequency	Number of core	Flops cycle (DP)	Peak Perf. (Gflops)
Xeon E5645 (Westmere)	2.4 GHz	2x6	4	115
Xeon E5-2687W0 (S.Bridge)	3.1 GHz	2x8	8	396
Xeon E5-2670v2 (I. Bridge)	2.5 GHz	2x10	8	400
Xeon E5-2630v3 (Hashwell)	2.4 GHz	2x8	16 (AVX-256bit)	614
Xeon E5-2697v4 (Broadwell)	2.3 GHz	2x18	16 (AVX-256bit)	1325
Xeon Platinum (Skylake)	2.1 GHz	2x24	32 (AVX-512bit)	3225

No increase	Factor 4x	Factor 8x	Total: Factor 32x

Real performance: serial

Performance ordered with respect to computational intensity (AI) =#flops/#byte
- CI > 1 FLOPs limited
- CI < 1 BW limited

GFLOP vs Computational Intensity (single core)

Real performance Improvements

Performance for single core performance

- OpenFoam, 3D lid-driven cavity, 80^3 grid-point, serial
- Linpack
- Stream
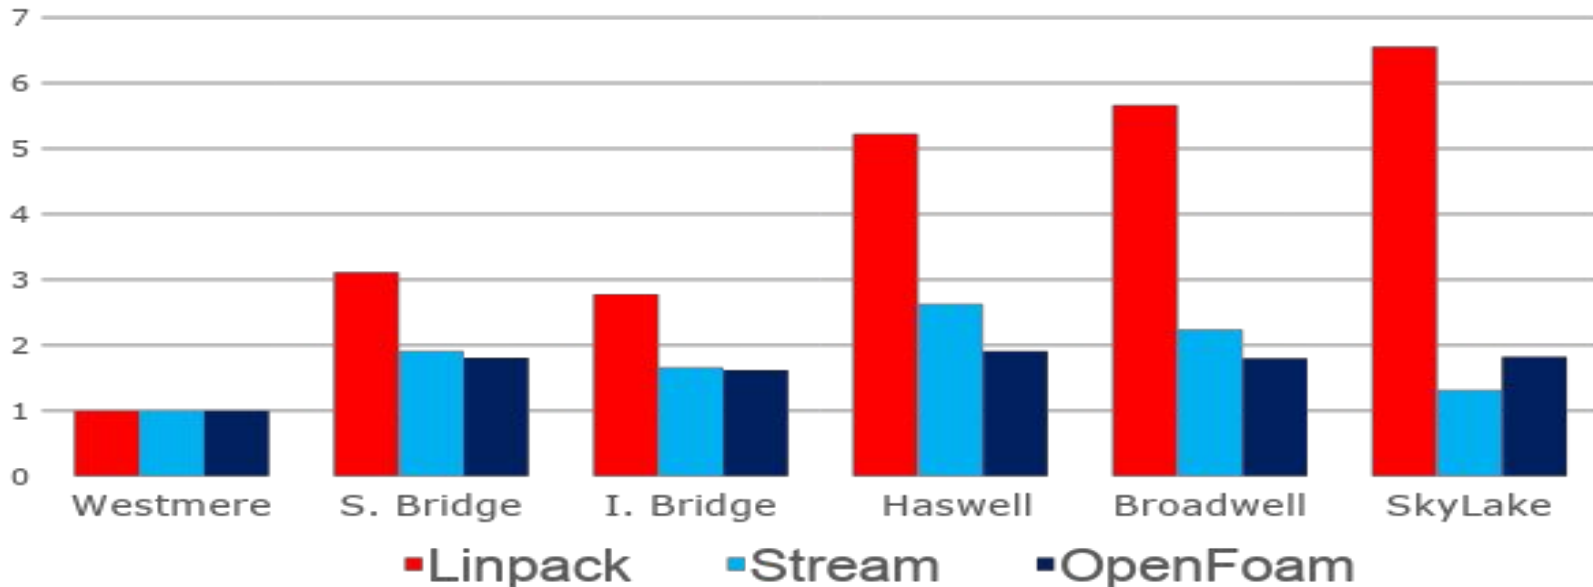
Real performance: parallel

Performance ordered with respect to computational intensity (AI): #flops/#byte

- CI > 1 FLOPs limited
- CI < 1 BW limited

GFLOP vs Computational Intensity (node)

Implementation

Implementation is the execution or practice of a plan, a method or any design, idea, model, specification, standard or policy for doing something.

✓ Starting from a well defined Algorithm the programmer write the program
- Programming languages
- data allocation
  - data movement
  - I/O
- compiler, tools
- HW

✓ Each of these points can improve/depress performance: are implementation issues

Instrument your code

Hint

✓ Instrument your code with timing functions
✓ Not all function but "logical" block (e.g. I/O, FFT, init, diagnostic)

```
call SYSTEM_CLOCK(countD0, count_rate, count_max)
// do something //
call SYSTEM_CLOCK(countD1, count_rate, count_max)
time_loop = real(countD1-countD0)/(count_rate)
```

```
#include <time.h>
…
time1 = clock();
// do something //
time2 = clock();
dub_time = (time2 - time1)/(double) CLOCKS_PER_SEC;
```

Tools

Many tool are (or were) used to (also) to profile/monitoring a code and much more, for example

✓ GNU perf
✓ GNU gprof
✓ Intel vtune
✓ Nvidia nsight
✓ AMD μprof
✓ scalasca
✓ likwid
✓ ….

gprof

✓ Old but useful: gprof flat profile

....
```
Flat profile:

Each sample counts as 0.01 seconds.
  %   cumulative   self              self     total
 time   seconds   seconds   calls  ms/call  ms/call  name
 34.25   0.13      0.13    3628800   0.00     0.00    check1
 31.61   0.25      0.12   36288000   0.00     0.00    d
 10.54   0.29      0.04          1  40.04    40.04    getgeom
  9.22   0.33      0.04   12470600   0.00     0.00    swap
  7.90   0.36      0.03   72576000   0.00     0.00    sqr
  5.27   0.38      0.02          1  20.02   340.34    search1
  1.32   0.38      0.01    3628800   0.00     0.00    save
  0.00   0.38      0.00          1   0.00   340.34    solve1
```

gprof

✓ Sld but useful: gprof call profile

```
index    %    time self   children   called    name
                                               <spontaneous>
[1] 100.0     0.00 0.38                         main [1]
              0.00 0.34      1/1                solve1 [3]
              0.04 0.00      1/1                getgeom [6]
-----------------------------------------------------
                            6235300             search1 [2]
              0.02 0.32        1/1              solve1 [3]
[2]  89.5     0.02 0.32    1+6235300            search1 [2]
              0.13 0.16  3628800/3628800        check1 [4]
              0.04 0.00 12470600/12470600       swap [7]
                            6235300             search1 [2]
-----------------------------------------------------
```

`gprof` **sintax**

A three step procedure

1. compile (all) the code with option "`-pg`"
   - `gcc -pg myfile.c -o myfile.x`
2. run the exe: it will create a gmon.out binary file
3. process the gmon.out file (standard output)
   - `gprof ./myfile.x > gprof.txt`

Caveat

✓ I could be intrusive
✓ Also linking step with flag "`-pg`"
✓ (unfortunately) is no more supported by some compilers :-(

Profile: Best Practices

✓ Check with your code complexity which performance are you expecting
- If they don't agree there's an issue

✓ Verify to be big enough not to fit entirely in cache
- Verify performance for all code/function/subroutines
- Cache hide real cade latencies

✓ Take care of profiling intrusivity
- Always take care of timing with or without profiling

✓ Use different Hw/Sw
- Can help to detecn "strange" behaviour

✓ Useful to understand the code flow
- especially if yuo don't have written the code!

MMM: implementation issues

Performance in Mflops (higher is better)

- ✓ HW: Intel(R) Xeon(R) Platinum 8260 CPU @ 2.40GHz
- ✓ Compiler intel 2021.5
    - ■ `ifort -O1`
- ✓ Matrix size: 2048 (96 MB)

Optimization	Mflops	Ratio
simple	2210	1
blocking (32)	2928	1.32x
unrolling x 4 (j)	3267	1.48x
unrolling x 4 (i)	2530	1.14x
unrolling x 2 (k)	3420	1.55x
All together	5653	2.56x

Example: Travel Salesman

https://en.wikipedia.org/wiki/Travelling_salesman_problem

Given a list of cities and the distances between each pair of cities, what is the shortest possible route that visits each city exactly once and returns to the origin city?

It is NP-complex problem: our approach

1. Brute force
2. Reducing combinations
3. Reducing functions
4. Reducing comparison
5. Precomputing distances

Brute-force approach: compute al N! combinations, where N is the number of cities to visit, unfortunately

- ✓ 10! = 3628800
- ✓ 11! = 479001600
- ✓ …
- ✓ 20!= 2432902008176640000

Really unfeasible! Any Idea?

```
time      seconds    seconds  calls   ms/call   ms/call   name
 34.25    0.13       0.13     3628800    0.00      0.00     check1
 31.61    0.25       0.12    36288000    0.00      0.00     d
 10.54    0.29       0.04           1   40.04     40.04     getgeom
  9.22    0.33       0.04    12470600    0.00      0.00     swap
  7.90    0.36       0.03    72576000    0.00      0.00     sqr
  5.27    0.38       0.02           1   20.02    340.34     search1
…
```

Reducing combinations

✓ We don't need to compute ALL the distancies. A lot are computed many times
✓ the distances with 4 cities (A,B,C,D)
  ■ ABCD = BCDA = CDAB = DABC
✓ You don't have to compute N! distances but (only) (n+1)! starting from a fixed city
✓ Big gain as N increase

```
time      seconds    seconds  calls   ms/call  ms/call  name
 66.73     0.02       0.02     362880   0.00     0.00     check1
 33.37     0.03       0.01    3628800   0.00     0.00     d
  0.00     0.03       0.00    7257600   0.00     0.00     sqr
  0.00     0.03       0.00    1247058   0.00     0.00     swap
  0.00     0.03       0.00     362880   0.00     0.00     save
```

Reducing functions

✓ We don't need to recompute all the distances for every single tour.
✓ Fix the distance between m cities and compute distance for remaining n-m cities

```
time     seconds    seconds       calls   ms/call   ms/call   name
35.75       0.03       0.03     1349289      0.00      0.00   d
21.45       0.04       0.02           1     15.02     65.07   search3
14.30       0.05       0.01     1247058      0.00      0.00   swap
14.30       0.06       0.01      362880      0.00      0.00   check3
 7.15       0.07       0.01     2698578      0.00      0.00   sqr
```

Check distances

✓ instead of check the final distance of a tour for the minimal stop the search if the
   (partial) tour is longer of the minimum found

```
time     seconds     seconds  calls   ms/call   ms/call   name
100.10    0.01        0.01    111515    0.00      0.00    d
  0.00    0.01        0.00    223030    0.00      0.00    sqr
  0.00    0.01        0.00    213374    0.00      0.00    swap
  0.00    0.01        0.00      2414    0.00      0.00    check3
  0.00    0.01        0.00      2414    0.00      0.00    save
```

Precomputing distancies

✓ The distance between A and B will be computed many time (i.e. build a **look-up table**)
✓ Precompute once and then access (matrix of n^2 elements)

```
time     seconds    seconds  calls   ms/call  ms/call  name
100.10    0.01       0.01          1   10.01    10.01   search5
  0.00    0.01       0.00     213374    0.00     0.00   swap
  0.00    0.01       0.00       2414    0.00     0.00   check5
  0.00    0.01       0.00       2414    0.00     0.00   save
  0.00    0.01       0.00        200    0.00     0.00   sqr
```

Travel Salesman: some figures

Results, in seconds, using IBM Power4@1100 Mhz (many time ago)

	9	10	11	12	13	14	15
1	0.92	10.3	123	1538	-	-	-
2	0.10	1.03	11.2	134	-	-	-
3	-	0.46	4.59	50.5	606	-	-
4	-	-	0.29	1.50	11.3	98.7	-
5	-	-	0.11	0.57	4.29	37.6	288

✓ Reducing unnecessary operation
✓ Removing repetition
✓ profiling help to focus on function to remove!

Travel Salesman: some (new) figures

Results, in seconds, using the AMD Ryzen 5 5625U (with gcc)

	13	14	15	16
4	0.66"	5.75"	43.6"	-
5	0.39"	3.33"	24.6"	170"
5 - initial condition	-	-	-	24.7"

- ✓ Reducing unnecessary operation
- ✓ Removing repetition
- ✓ a good initial condition can help (from rel. 4)
- ✓ profiling help to focus on operations to remove/optimize!
- ✓ HW improvement helps, but SW are bigger!!!

Travel Salesman: 16 cities

Sieve of Eratosthenes

✓ https://en.wikipedia.org/wiki/Sieve_of_Eratosthenes
✓ One of the first known Algorithm for finding prime numbers
  - much older than computers!
  - It is an iterative algorithm

	2	3	4	5	6	7	8	9	10	Prime numbers
11	12	13	14	15	16	17	18	19	20	
21	22	23	24	25	26	27	28	29	30	
31	32	33	34	35	36	37	38	39	40	
41	42	43	44	45	46	47	48	49	50	
51	52	53	54	55	56	57	58	59	60	
61	62	63	64	65	66	67	68	69	70	
71	72	73	74	75	76	77	78	79	80	
81	82	83	84	85	86	87	88	89	90	
91	92	93	94	95	96	97	98	99	100	
101	102	103	104	105	106	107	108	109	110	
111	112	113	114	115	116	117	118	119	120	

Sieve of Eratosthenes: step0

✓ https://en.wikipedia.org/wiki/Sieve_of_Eratosthenes

✓ First implementation

```fortran
parameter(nd=50000)                   ! max number to check
…
do i = 2, nd
   a(i) = 0                           ! possible prime number
enddo
…
do icheck = 2, nd
    do i = icheck+1, nd
        if (mod(i,icheck).eq.0) then
            a(i) = 1                  ! i is not a prime number
        endif
    enddo
enddo
```

Sieve of Eratosthenes: figures

✓ Intel PentiumIII, 700 Mhz   (old stuff)

	Time (s)	Gain (%)
Step0	5118	-

Sieve of Eratosthenes: step 1

✓ The original code performs a lot of "useless" check.
✓ It checks, for example, all the number that are divisible by 4, that are a subset of those divisible by 2.
✓ You have to check only with numbers that you know are candidate as prime, i.e. those with a(i)=0

```fortran
parameter(nd=50000)                    ! max number to check
…
do icheck = 2, nd
   do i = icheck+1, nd
      if (a(i).eq.0) then
          if (mod(i,icheck).eq.0) then
              a(i) = 1                  ! i is not a prime number
          endif
      endif
   enddo
enddo
```

Sieve of Eratosthenes: figures

✓ Intel PentiumIII, 700 Mhz   (old stuff)

	Time (s)	Gain (%)
Step0	5118	-
Step1	2383	53.5%

✓ Other "useless" operation to skip

✓ If you know that a number is not a prime you can skip the check

```fortran
parameter(nd=50000)                    ! max number to check
…
do icheck = 2, nd
    if (a(icheck).eq.0) then
        do i = icheck+1, nd
            if (a(i).eq.0) then
                if (mod(i,icheck).eq.0) then
                    a(i) = 1           ! i is not a prime number
                endif
            endif
        enddo
    endif
enddo
```

Sieve of Eratosthenes: figures

✓ Intel PentiumIII, 700 Mhz   (old stuff)

	Time (s)	Gain (%)
Step0	5118	-
Step1	2383	53.5%
Step2	212	95.8%

Sieve of Eratosthenes: step 3

✓ Other "useless" operation to skip
✓ If m is the biggest prime number found you must start the search from m*m

```fortran
parameter(nd=50000)                    ! max number to check
…
do icheck = 2, nd
    if (a(icheck).eq.0) then
        do i = icheck*icheck, nd
            if (a(i).eq.0) then
                if (mod(i,icheck).eq.0) then
                    a(i) = 1           ! i is not a prime number
                endif
            endif
        enddo
    endif
enddo
```

Sieve of Eratosthenes: figures

✓   Intel PentiumIII, 700 Mhz   (old stuff)

	Time (s)	Gain (%)
Step0	5118	-
Step1	2383	53.5%
Step2	212	95.8%
Step3	1.21	99.976%

✓ Instead of check modulo (i.e. a division) we mark directly the multiple (i.e. multiplication)

```fortran
parameter(nd=50000)               ! max number to check
…
do icheck = 2, nd
    if (a(icheck).eq.0) then
        imax = INT(nd/i)
        do i=1,imax
            a(i*icheck) = 1       ! i is not a prime number
        enddo
    endif
enddo
```

Sieve of Eratosthenes: figures

✓ Intel PentiumIII, 700 Mhz   (old stuff)

	Time (s)	Gain (%)
Step0	5118	-
Step1	2383	53.5%
Step2	212	95.8%
Step3	1.21	99.976%
Step4	0.13	99.9974%

Sieve of Eratosthenes: step 5 (fine tuning)

✓ Mark as non-prime number up to 30 in the initialization step
✓ Where is the problem?

```
a( 2) = 0                     ! prime
a( 3) = 0                     ! prime
a( 4) = 1
a( 5) = 0                     ! prime
....
a(30) = 1
do i=30,nd,30
   a(i+ 1) = 0                ! could be a prime
   a(i+ 2) = 1
   a(i+ 3) = 1
   a(i+ 4) = 1
   a(i+ 5) = 1
     ....
   a(i+29) = 0                 ! could be a prime
   a(i+30) = 1
enddo
```

Sieve of Eratosthenes: figures

✓ Intel PentiumIII, 700 Mhz (old stuff)

	Time (s)	Gain (%)
Step0	5118	-
Step1	2383	53.5%
Step2	212	95.8%
Step3	1.21	99.976%
Step4	0.13	99.9974%
Step5	0.094	99.9982%

Sieve of Eratosthenes: step 6 (fine tuning)

✓ Limit the search up to sqrt(nd)

    ■   it could be is the biggest prime number up to nd

```fortran
parameter(nd=50000)                     ! max number to check
…
do icheck = 7, INT(sqrt(float(nd))
   if (a(icheck).eq.0) then
      imax = INT(nd/i)
      do i=1,imax
         a(i*icheck) = 1              ! i is not a prime number
      enddo
   endif
enddo
```

Sieve of Eratosthenes: figures

✓ Intel PentiumIII, 700 Mhz   (old stuff)

	Time (s)	Gain (%)
Step0	5118	-
Step1	2383	53.5%
Step2	212	95.8%
Step3	1.21	99.976%
Step4	0.13	99.9974%
Step5	0.094	99.9982%
Step6	0.078	99.9985%

Who is in charge of what (reprise)?

High performance is the joint effort of different "players"

- ✓ **Programmer**: in charge of the choice of the algorithm
- ✓ **Compiler**: in charge of the traduction in instructions
- ✓ **HW**: in charge of execute the instructions

Different operations, different cost

✓ Different costs
  - sum → few cycles
  - product → few cycles
  - division → many cycles
  - square root → many many cycles
  - exponent → many cycles
  - trigonometric functions → many cycles
  - data access → few/many access

✓ (possible) Solutions
  - Lookup table
  - If range limited
    - polynomial approximation
    - taylor series
    - prostaferesi formula

✓ Take care of the chosen Algorithm
   - You (should) know how it works and which operations perform or remove (e.g. TSP)
✓ Take care of Algorithm Implementation
   - Data allocation
   - Data access
   - Operation to be performed
✓ timing/profiling could help
   - in finding bottlenecks
   - in understanding where you're wrong
✓ In general: coding style could seriously affect performance
   - Loop ordering
✓ Some operation are more "expensive" with respect to others

(some) Keywords

- ✓ **Look-up table**: precomputed data. It could be a gain in time (size dependent)
- ✓ **Roofline model**
- ✓ **Algorithm complexity**
- ✓ **Bandwidth limited code**
- ✓ **FP limited code**

Some references

✓ Bentley, More Programming Pearls (2nd Ed.), ACM Press/Addison-Wesley, 1989

✓ Bentley, DDJ, 1999

✓ Discovering faster matrix multiplication algorithms with reinforcement learning, Nature, 2023,

✓ Dispense Master HPC 2007